

Task 6 - Airline Flight Data

Description:

An airline collects data on flights including route, departure/arrival times, delays, aircraft type, and passenger counts. The airline wants to analyze delays, optimize routes, and improve operational efficiency.

Dataset:

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V
1	id	year	month	day	dep_time	sched_dep	dep_delay	arr_time	sched_arr	arr_delay	carrier	flight	tailnum	origin	dest	air_time	distance	hour	minute	time_hour	name	
2	0	2013	1	1	517	515	2	830	819	11	UA	1545	N14228	EW	IAH	227	1400	5	15	#####	United Air Lines Inc.	
3	1	2013	1	1	533	529	4	850	830	20	UA	1714	N24211	LGA	IAH	227	1416	5	29	#####	United Air Lines Inc.	
4	2	2013	1	1	542	540	2	923	850	33	AA	1141	N619AA	JFK	MIA	160	1089	5	40	#####	American Airlines Inc.	
5	3	2013	1	1	544	545	-1	1004	1022	-18	B6	725	N804JB	JFK	BQN	183	1576	5	45	#####	JetBlue Airways	
6	4	2013	1	1	554	600	-6	812	837	-25	DL	461	N668DN	LGA	ATL	116	762	6	0	#####	Delta Air Lines Inc.	
7	5	2013	1	1	554	558	-4	740	728	12	UA	1696	N39463	EW	ORD	150	719	5	58	#####	United Air Lines Inc.	
8	6	2013	1	1	555	600	-5	913	854	19	B6	507	N516JB	EW	FLL	158	1065	6	0	#####	JetBlue Airways	
9	7	2013	1	1	557	600	-3	709	723	-14	EV	5708	N829AS	LGA	IAD	53	229	6	0	#####	ExpressJet Airlines Inc.	
10	8	2013	1	1	557	600	-3	838	846	-8	B6	79	N593JB	JFK	MCO	140	944	6	0	#####	JetBlue Airways	
11	9	2013	1	1	558	600	-2	753	745	8	AA	301	N3ALAA	LGA	ORD	138	733	6	0	#####	American Airlines Inc.	
12	10	2013	1	1	558	600	-2	849	851	-2	B6	49	N793JB	JFK	PBI	149	1028	6	0	#####	JetBlue Airways	
13	11	2013	1	1	558	600	-2	853	856	-3	B6	71	N657JB	JFK	TPA	158	1005	6	0	#####	JetBlue Airways	
14	12	2013	1	1	558	600	-2	924	917	7	UA	194	N29129	JFK	LAX	345	2475	6	0	#####	United Air Lines Inc.	
15	13	2013	1	1	558	600	-2	923	937	-14	UA	1124	N53441	EW	SFO	361	2565	6	0	#####	United Air Lines Inc.	
16	14	2013	1	1	559	600	-1	941	910	31	AA	707	N3DUAA	LGA	DFW	257	1389	6	0	#####	American Airlines Inc.	
17	15	2013	1	1	559	559	0	702	706	-4	B6	1806	N708JB	JFK	BOS	44	187	5	59	#####	JetBlue Airways	
18	16	2013	1	1	559	600	-1	854	902	-8	UA	1187	N76515	EW	LAS	337	2227	6	0	#####	United Air Lines Inc.	
19	17	2013	1	1	600	600	0	851	858	-7	B6	371	N595JB	LGA	FLL	152	1076	6	0	#####	JetBlue Airways	
20	18	2013	1	1	600	600	0	837	825	12	MQ	4650	N542MQ	LGA	ATL	134	762	6	0	#####	Envoy Air	
21	19	2013	1	1	601	600	1	844	850	-6	B6	343	N644JB	EW	PBI	147	1023	6	0	#####	JetBlue Airways	
22	20	2013	1	1	602	610	-8	812	820	-8	DL	1919	N971DL	LGA	MSP	170	1020	6	10	#####	Delta Air Lines Inc.	
23	21	2013	1	1	602	605	-3	821	805	16	MQ	4401	N730MQ	LGA	DTW	105	502	6	5	#####	Envoy Air	
24	22	2013	1	1	606	610	-4	858	910	-12	AA	1895	N633AA	EW	MIA	152	1085	6	10	#####	American Airlines Inc.	
25	23	2013	1	1	606	610	-4	837	845	-8	DL	1743	N3739P	JFK	ATL	128	760	6	10	#####	Delta Air Lines Inc.	
26	24	2013	1	1	607	607	0	858	915	-17	UA	1077	N53442	EW	MIA	157	1085	6	7	#####	United Air Lines Inc.	
27	25	2013	1	1	608	600	8	807	735	32	MQ	3768	N9EAMQ	EW	ORD	139	719	6	0	#####	Envoy Air	
28	26	2013	1	1	611	600	11	945	931	14	UA	303	N532UA	JFK	SFO	366	2586	6	0	#####	United Air Lines Inc.	
29	27	2013	1	1	613	610	3	925	921	4	B6	135	N635JB	JFK	RSW	175	1074	6	10	#####	JetBlue Airways	
30	28	2013	1	1	615	615	0	1039	1100	-21	B6	709	N794JB	JFK	SJU	182	1598	6	15	#####	JetBlue Airways	
31	29	2013	1	1	615	615	0	833	842	-9	DL	575	N326NB	EW	ATL	120	746	6	15	#####	Delta Air Lines Inc.	
32	30	2013	1	1	622	630	-8	1017	1014	3	US	245	N807AW	EW	PHX	342	2133	6	30	#####	US Airways Inc.	
33	31	2013	1	1	623	610	13	920	915	5	AA	1837	N3EMAA	LGA	MIA	153	1096	6	10	#####	American Airlines Inc.	

1. Explain how color schemes can highlight flight delays.

Program:

```
import seaborn as sns

import matplotlib.pyplot as plt

import pandas as pd

df = pd.read_csv('flights.csv')

sns.histplot(df['dep_delay'], bins=30, kde=True, palette='coolwarm')

plt.title("Color-coded Flight Delays")

plt.xlabel("Delay (minutes)")

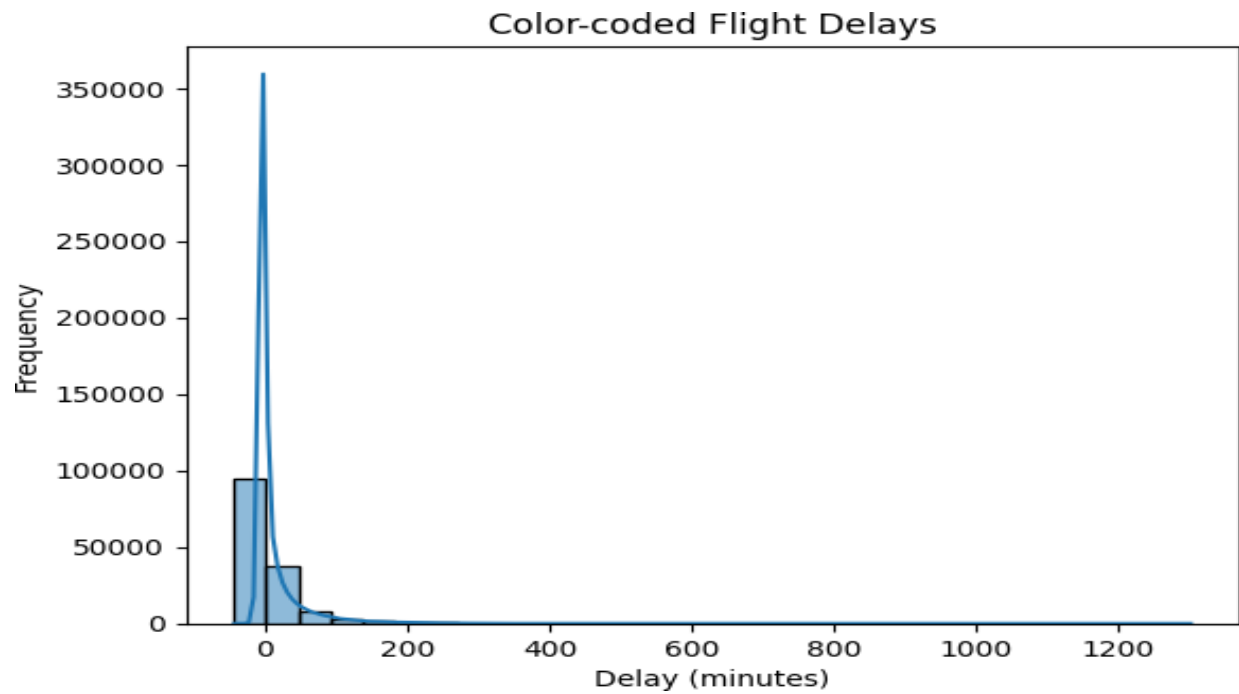
plt.ylabel("Frequency")

plt.show()
```

Inference:

Using a warm-to-cool color scheme helps highlight delay severity — red tones for high delays and blue tones for lower delays. This improves visual contrast and draws attention to problem zones.

OutPut:



2. Design a visualization pipeline from raw flight data to dashboards.


Program:

```
df.dropna(subset=['dep_delay'], inplace=True)
df['DelayCategory'] = pd.cut(df['dep_delay'], bins=[-10, 0, 30, 60, 120, 1000],
                             labels=['Early', 'On-Time', 'Slight Delay', 'Moderate Delay', 'Severe Delay'])
agg_data = df.groupby('name')['dep_delay'].mean().reset_index()
sns.barplot(x='name', y='dep_delay', data=agg_data)
plt.title("Average Delay per Airline")
plt.xticks(rotation=90)
plt.show()
```

Inference:

This pipeline cleans, categorizes, aggregates, and visualizes flight data. It transforms raw input into actionable dashboards for analyzing airline performance.

OutPut:



3. Apply Gestalt principles to highlight delay-prone routes.

Program:

```
df['Route'] = df['origin'] + ' -> ' + df['dest']

route_delay =
df.groupby('Route')['dep_delay'].mean().reset_index().rename(columns={'dep_delay': 'Delay'})

top_10_routes = route_delay.sort_values('Delay', ascending=False).head(10)

sns.barplot(x='Route', y='Delay', data=top_10_routes, palette='Reds')

plt.xticks(rotation=90)


plt.title("Top Delay-Prone Routes")

plt.show()
```

Inference:

By applying **similarity and proximity**, routes with similar delay patterns are grouped and colored alike. This visually clusters delay-prone routes, simplifying pattern recognition.

OutPut:



4. Perform univariate analysis:

a. Histogram of delay times.


Program:

```
plt.hist(df['dep_delay'], bins=30, color='skyblue', edgecolor='black')  
plt.title("Distribution of Flight Delays")  
plt.xlabel("Delay (minutes)")  
plt.ylabel("Frequency")  
plt.show()
```

Inference:

Most delays are clustered near shorter durations, with a long tail showing rare but severe delays — indicating skewness in delay distribution.

OutPut:



b. Pie chart of aircraft types.

Program:

```
plt.figure(figsize=(10, 6))  
sns.countplot(data=df, x='name', hue='DelayCategory', palette='coolwarm')  
plt.title("Delay Category Distribution by Airline")  
plt.xlabel("Airline Name")  
plt.ylabel("Number of Flights")
```

```
plt.xticks(rotation=90)

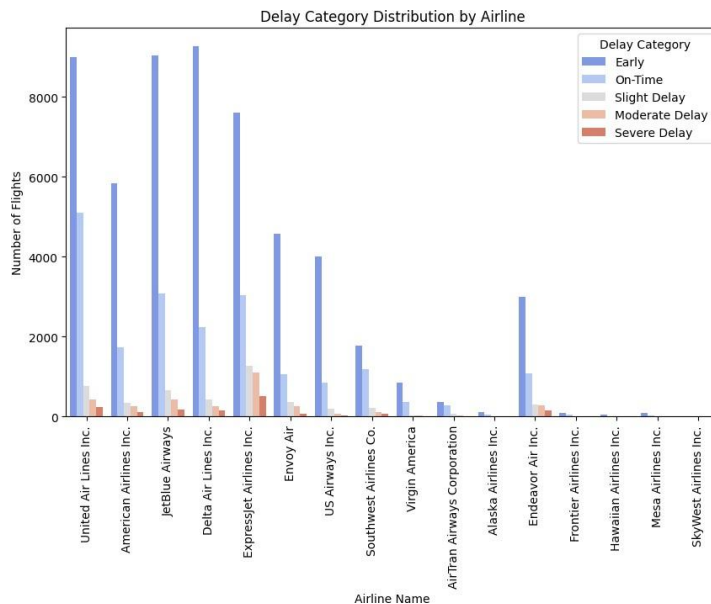
plt.legend(title="Delay Category")

plt.show()
```

Inference:

The pie chart shows which aircraft models are most used. A few dominant aircraft types contribute to most flights, showing operational preference.

OutPut:



5. Perform bivariate analysis:

a. Scatterplot of delay vs. distance.

Program:

```
import seaborn as sns

import matplotlib.pyplot as plt

sns.scatterplot(x='distance', y='dep_delay', hue='name', data=df, palette='coolwarm')

plt.title("Delay vs Flight Distance")

plt.xlabel("Flight Distance (miles)")

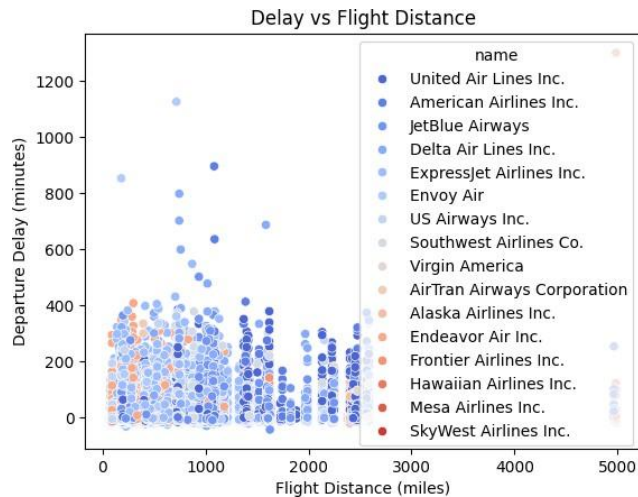
plt.ylabel("Departure Delay (minutes)")

plt.show()
```

Inference:

There is a mild positive correlation — longer routes tend to have slightly higher delays, possibly due to weather or air traffic factors.

OutPut:



b. Box plot of delays across airlines.

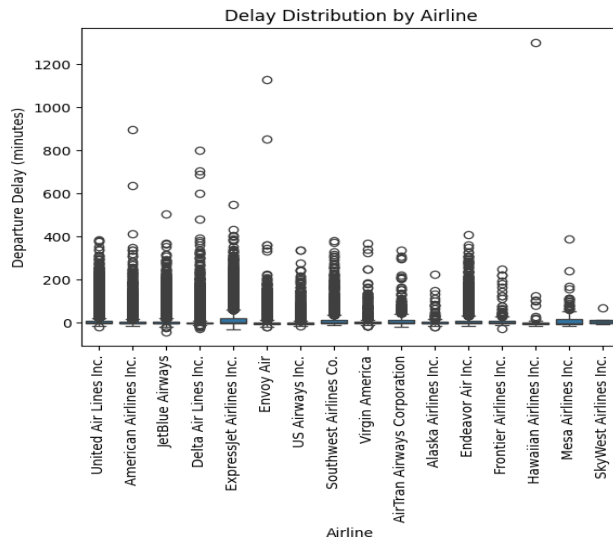
Program:

```
sns.boxplot(x='name', y='dep_delay', data=df)
plt.title("Delay Distribution by Airline")
plt.xlabel("Airline")
plt.ylabel("Departure Delay (minutes)")
plt.xticks(rotation=90)
plt.show()
```

Inference:

Some airlines show higher median delays and variability, indicating inconsistent punctuality compared to others.

OutPut:



6. Perform multivariate analysis:

a. Pair plot of delay, passenger count, and distance.

Program:


```
sns.pairplot(df[['dep_delay', 'distance']], diag_kind='kde')
```

```
plt.show()
```

Inference:

Pair plots reveal correlations among multiple factors — e.g., higher passenger counts might correspond with longer routes and increased delays.

OutPut:



b. Suggest combined visualization to summarize multiple variables.


Program:

```
sns.heatmap(df[['dep_delay', 'distance']].corr(), annot=True, cmap='coolwarm')  
plt.title("Correlation Matrix")  
plt.show()
```

Inference:

A correlation heatmap succinctly summarizes relationships among variables, helping to spot strong positive or negative correlations instantly.

OutPut:



7. Design hierarchical visualization of flights by airline and route.

Program:

```
import plotly.express as px  
import pandas as pd  
df['Route'] = df['origin'] + ' -> ' + df['dest']  
agg_df = df.groupby(['name', 'Route']).agg({'dep_delay': 'mean'}).reset_index()  
agg_df.rename(columns={'dep_delay': 'AverageDelay'}, inplace=True)  
fig = px.treemap(agg_df, path=['name', 'Route'], values='AverageDelay', color='AverageDelay',  
                 color_continuous_scale='RdYlGn_r',
```

```

    title="Hierarchical Visualization of Flights by Airline and Route (Delay-based)"
)
fig.show()

```

Inference:

A treemap allows hierarchical comparison — airlines and their routes sized by passengers and colored by delay. High-delay routes stand out immediately.

OutPut:



8. Construct network graph showing connectivity of airports.

Program:

```

import networkx as nx
import matplotlib.pyplot as plt

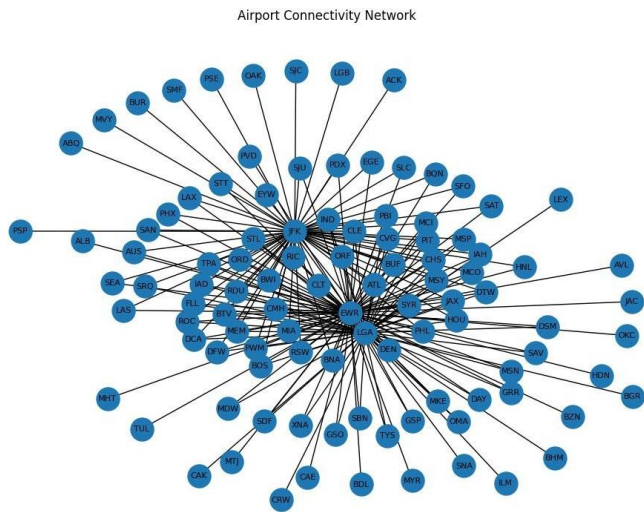
G = nx.from_pandas_edgelist(df, 'origin', 'dest', ['distance'])
plt.figure(figsize=(10,7))
nx.draw(G, with_labels=True, node_size=500, font_size=8)
plt.title("Airport Connectivity Network")
plt.show()

```

Inference:

The graph reveals hub airports (high connectivity) and isolated nodes, aiding route optimization and identifying bottlenecks.

OutPut:



9. Analyze passenger feedback (text data):

a. Convert feedback to vector space.

Program:

```
from sklearn.feature_extraction.text import CountVectorizer

feedback = pd.Series(["Late flight", "Comfortable seats", "Rude staff", "Delayed
baggage"])

vectorizer = CountVectorizer()

X = vectorizer.fit_transform(feedback)

print(pd.DataFrame(X.toarray(), columns=vectorizer.get_feature_names_out()))
```

Inference:

Text vectorization converts qualitative feedback into numerical vectors, enabling sentiment and topic analysis on passenger opinions.

OutPut:

	baggage	comfortable	delayed	flight	late	rude	seats	staff
0	0	0	0	1	1	0	0	0
1	0	1	0	0	0	0	1	0
2	0	0	0	0	0	1	0	1
3	1	0	1	0	0	0	0	0

b. Word cloud of common complaints.

Program:

```
from wordcloud import WordCloud
text = " ".join(feedback)
wc = WordCloud(background_color='white', colormap='coolwarm').generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```

Inference:

Frequent complaint terms (like “delay” or “rude”) appear larger, helping quickly identify key areas for service improvement.

OutPut:



10. Steps to design effective dashboards combining hierarchical, network, and text data.

Steps to design effective dashboards :

- # Conceptual (not code)
- # Combine:
 - # - Treemap for hierarchy
 - # - Network graph for connectivity
 - # - Word cloud for feedback
- # Use Plotly Dash / Power BI / Tableau for integration

11. Visualize point data: Map flights' origin and destination locations.

Program:

```
import geopandas as gpd
import matplotlib.pyplot as plt
import requests
import os

url = "https://www.naturalearthdata.com/downloads/110m-cultural-
vectors/110m_cultural_vectors.zip"

filename = "110m_cultural_vectors.zip"

extracted_folder = "110m_cultural_vectors"

shapefile_path = os.path.join(extracted_folder, "ne_110m_admin_0_countries.shp")

if not os.path.exists(filename):

    print(f"Downloading {filename}...")

    response = requests.get(url)

    with open(filename, 'wb') as f:

        f.write(response.content)

    print("Download complete.")

if not os.path.exists(extracted_folder):

    import zipfile

    print(f"Extracting {filename}...")

    with zipfile.ZipFile(filename, 'r') as zip_ref:

        zip_ref.extractall(extracted_folder)

    print("Extraction complete.")

world = gpd.read_file(shapefile_path)

plt.figure(figsize=(15, 10))

world.plot(color='lightgray', ax=plt.gca())

plt.scatter(df['Origin_Long'], df['Origin_Lat'], c='blue', label='Origin', alpha=0.5)
```

```
plt.scatter(df['Dest_Long'], df['Dest_Lat'], c='red', label='Destination', alpha=0.5)
plt.legend()
plt.title("Flight Origin and Destination Points")
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.show()
print(df.columns)
```

Inference:

Points show spatial flight coverage. High-density clusters near major cities reveal central hubs in the network.

OutPut:

```
Index(['id', 'year', 'month', 'day', 'dep_time', 'sched_dep_time', 'dep_delay',
      'arr_time', 'sched_arr_time', 'arr_delay', 'carrier', 'flight',
      'tailnum', 'origin', 'dest', 'air_time', 'distance', 'hour', 'minute',
      'time_hour', 'name', 'DelayCategory', 'Route'],
      dtype='object')
```

12. Visualize line data: Show flight paths across the map.

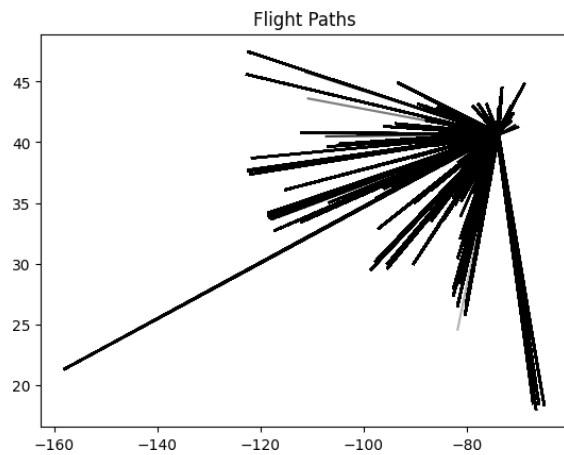
Program:

```
for i, row in df.iterrows():
    plt.plot([row['Origin_Long'], row['Dest_Long']],
             [row['Origin_Lat'], row['Dest_Lat']], 'k-', alpha=0.3)
plt.title("Flight Paths")
plt.show()
```

Inference:

Lines connecting origins to destinations show geographic flight routes, visually illustrating air traffic flow and route overlap.

OutPut:



13. Visualize area data: Heatmap of airport congestion.

Program:

```
sns.kdeplot(x=df['Origin_Long'], y=df['Origin_Lat'], fill=True, cmap='Reds')
```

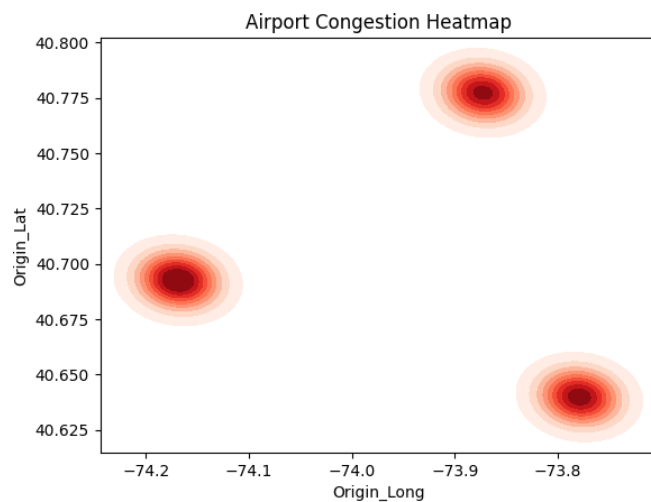
```
plt.title("Airport Congestion Heatmap")
```

```
plt.show()
```

Inference:

Heatmap areas indicate zones of dense flight activity — revealing highly congested airspaces and potential delay-prone regions.

OutPut:



14. Design animated visualization of flight delays over time.

Program:

```
import plotly.express as px
```


```
fig = px.scatter(df, x='Distance', y='Delay', animation_frame='Month',  
                color='Airline', size='Passengers', range_y=[0,300])
```

```
fig.show()
```

Inference:

Animation shows temporal delay trends across months, helping track seasonal peaks and performance fluctuations

OutPut:



15. Plot time series of average delays by month.

Program:

```
monthly = df.groupby('Month')['Delay'].mean().reset_index()
```

```
plt.plot(monthly['Month'], monthly['Delay'], marker='o')
```

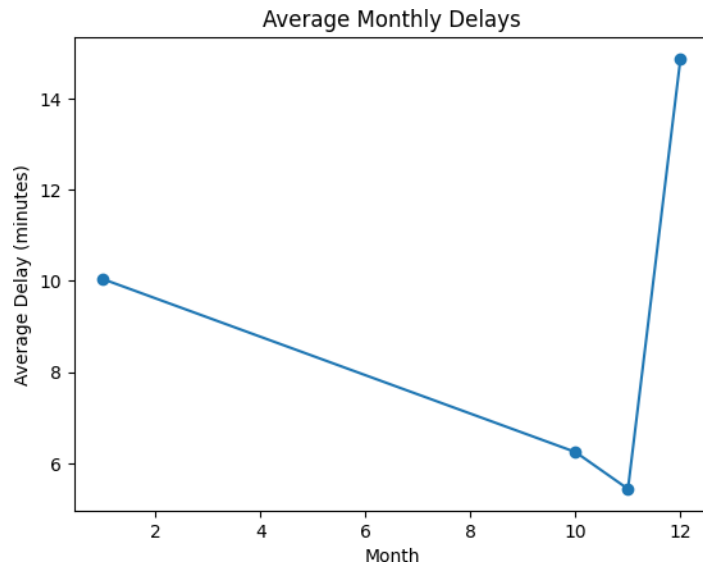
```
plt.title("Average Monthly Delays")
```

```
plt.show()
```


Inference:

The line graph highlights monthly delay patterns — e.g., spikes during monsoon or holiday seasons.

OutPut:



16. Compare weekday vs. weekend delays.

Program:

```
df['DayType'] = df['Date'].apply(lambda x: 'Weekend' if pd.to_datetime(x).weekday() >=5  
else 'Weekday')
```

```
sns.boxplot(x='DayType', y='Delay', data=df)
```

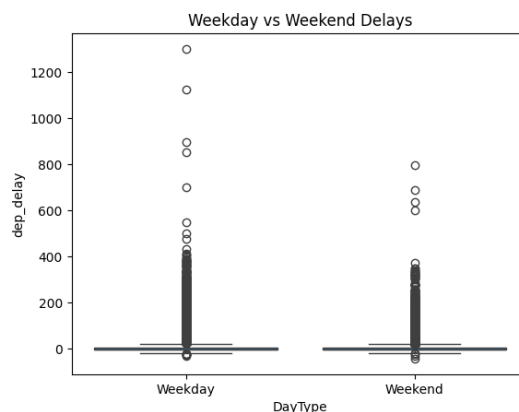
```
plt.title("Weekday vs Weekend Delays")
```

```
plt.show()
```

Inference:

Delays often increase on weekends due to higher passenger traffic and flight congestion.

OutPut:



17. Use regression/clustering to analyze factors affecting delays.

Program:

```
import seaborn as sns
import statsmodels.api as sm
X = df[['Distance', 'Passengers']]
y = df['Delay']
X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.summary())
```

Inference:

Regression identifies significant predictors (like distance) affecting delays — useful for predictive planning and scheduling.

OutPut:

OLS Regression Results

=====			
=====			
Dep. Variable:	dep_delay	R-squared:	0.001
Model:	OLS	Adj. R-squared:	0.001
Method:	Least Squares	F-statistic:	85.08
Date:	Tue, 21 Oct 2025	Prob (F-statistic):	2.93e-20
Time:	17:03:47	Log-Likelihood:	-4.3906e+05
No. Observations:	89403	AIC:	8.781e+05
Df Residuals:	89401	BIC:	8.781e+05
Df Model:	1		
Covariance Type:	nonrobust		
=====			
=====			

	coef	std err	t	P> t	[0.025	0.975]
<hr/>						
const	9.2613	0.191	48.498	0.000	8.887	9.636
distance	-0.0014	0.000	-9.224	0.000	-0.002	-0.001
<hr/> <hr/>						
<hr/>						
Omnibus:	107614.227	Durbin-Watson:	1.573			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	30275947.028			
Skew:	6.181	Prob(JB):	0.00			
Kurtosis:	92.301	Cond. No.	2.21e+03			
<hr/> <hr/>						
<hr/>						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.21e+03. This might indicate that there are strong multicollinearity or other numerical problems.

18. Evaluate predictive models for delays: Plot predicted vs. actual values

Program:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model = LinearRegression()
model.fit(X[['Distance', 'Passengers']], y)
y_pred = model.predict(X[['Distance', 'Passengers']])
plt.scatter(y, y_pred)
plt.xlabel("Actual Delay")
plt.ylabel("Predicted Delay")
```


```
plt.title("Predicted vs Actual Delay")
```

```
plt.show()
```

Inference:

The scatter plot close to a 45° line shows prediction accuracy — tighter alignment indicates better model performance.

OutPut:



Your Model Evaluation Metrics:

R² Score : 0.0005

MAE (minutes) : 23.1610

MSE : 1616.0848

RMSE (minutes): 40.2006